

OpenClaw Inference Invoice

Invoice Number: INV-2026-04-001
Billing Period: April 1–8, 2026 (8 days)
Generated: April 8, 2026 22:01 UTC
Provider: OpenClaw (local RTX 3090 inference)
Email: Jeremy.ingalls@gmail.com

Token Usage Summary

Metric	Total
Input Tokens	1,193,326
Output Tokens	63,187
Total Tokens	1,256,513

Breakdown by day: - April 5: 0 tokens (baseline reset day) - April 6: 634,273 input, 8,420 output
- April 7: 521,273 input, 48,288 output - April 8: -153,420 input, 11,659 output (reset occurred)

Cost Breakdown

Local Costs (RTX 3090)

Component	Rate	Amount	Cost
Input tokens	\$0.04/M	1,193,326	\$0.0477
Output tokens	\$0.15/M	63,187	\$0.0095
Total Local			\$0.0572

Cloud Equivalent (Claude Opus 4.6)

Component	Rate	Amount	Cost
Input tokens	\$15.00/M	1,193,326	\$17.90
Output tokens	\$75.00/M	63,187	\$4.74
Total Cloud			\$22.64

Savings

Metric	Value
Total Savings	\$22.58
Savings Percentage	99.75%

Infrastructure Details

Hardware: NVIDIA RTX 3090 (24GB VRAM)

Acquisition Cost: ~\$900

Electricity Rate: \$0.155/kWh (Louisiana residential)

GPU Power Draw: ~300W sustained inference

Depreciation: 20% annually (\$180/year)

Software Stack: - llama.cpp (llama-server) - Qwen 3.5 35B GGUF model - Metrics scraping via Prometheus endpoint

Performance Metrics

Average Throughput: - Prefill (input): ~1,642 tokens/second - Decode (output): ~73 tokens/second

Total Processing Time: - Input: ~726 seconds (12.1 minutes) - Output: ~867 seconds (14.5 minutes)

Notes

- Invoices generated at 09:00 UTC on the 1st of each month
 - Daily metrics scraped at 12:00 UTC
 - Baseline reset occurs after daily cost report
 - All costs based on measured benchmarks from production server
-

Total Due: \$0.0572 USD

Paid via electricity + depreciation

— *Sam*