

## Ollama Model Cost Invoice – April 9, 2026 (2 P6000 GPUs)

Model: llama/llamacpp (running on 2 NVIDIA P6000 GPUs @ 250 W each, total 500 W)

Assumptions: - Power draw: 500 W (0.5 kW) sustained inference - Electricity rate: \$0.10 per kWh (US residential average) - Hours of operation today: 6 h (estimate) - Tokens processed today: 1,068,832 tokens (0.8 M input + 0.067 M output – same as the 3090 example for a fair comparison) - Acquisition cost of each P6000: \$1500 (total \$3000) - Depreciation: 20 % annually -> \$600 / yr ~ \$1.64 / day

Cost Breakdown (local hardware):

| Component    | Rate       | Usage              | Amount     | Cost                                         |
|--------------|------------|--------------------|------------|----------------------------------------------|
| Input tokens | \$0.04 / M | tokens 1,002,126   |            |                                              |
| 1.002126 M   | \$0.0401   | Output tokens      | \$0.15 / M | tokens 66,706                                |
| 0.066706 M   | \$0.0100   | Token cost (local) | \$0.0501   | Electricity (0.5 kW x 6 h x \$0.10/kWh)      |
|              |            |                    | \$0.30     | Depreciation ( \$600 / yr ÷ 365 days × 6 h ) |
|              |            |                    | \$0.13     | Total local cost                             |
|              |            |                    | \$0.48     |                                              |

Cloud-equivalent (Claude Opus 4.6) for comparison: - Input: \$15.00 / M x 1.002 M = \$15.03 - Output: \$75.00 / M x 0.067 M = \$5.00 Total cloud cost: \$20.03 Savings vs. cloud: \$19.55 (~ 97 %)

Infrastructure details - Hardware: 2 x NVIDIA P6000, 24 GB VRAM each - Power draw: ~250 W per GPU, 500 W total - Electricity rate: \$0.10/kWh (Louisiana residential – same as example) - Software stack: llama.cpp (llama-server) with Qwen 3.5 35B GGUF model - Metrics scraped via Prometheus at 12:00 UTC; daily baseline reset

Generated automatically by Sam, your ops-butler.