

Ollama Model Cost Invoice – April 9, 2026

Customer: Jeremy Ingalls **Provider:** Sam – Operations Butler **Date:** April 9, 2026

Usage Summary

- Model: llama/llamacpp (Qwen 3.5 35B GGUF)
- Hardware: 2 x NVIDIA P6000 GPUs, 250 W each (total 500 W)
- Tokens processed: 1,068,832 (1,002,126 input, 66,706 output)

Cost Breakdown

Component	Rate	Amount	Cost
Input tokens	\$0.04 per M tokens	1,002,126	\$0.0401
Output tokens	\$0.15 per M tokens	66,706	\$0.0100
Token cost (local)			\$0.0501
Electricity (0.5 kW x 6 h x \$0.16/kWh)			\$0.48
Depreciation (20% yr on \$3,000 hardware)			\$0.13
Total local cost			\$0.66

Cloud Equivalent (Claude Opus 4.6)

- Input tokens: \$15.00 per M -> \$15.03
- Output tokens: \$75.00 per M -> \$5.00
- **Total cloud cost:** \$20.03
- **Savings:** \$19.37 (~ 97%)

Infrastructure Details

- Acquisition cost: \$3,000
- Depreciation: \$1.64 per day, \$0.13 for today's 6-hour window
- Electricity rate: \$0.16 per kWh (Louisiana residential)
- Monitoring: Prometheus, daily reset at 12:00 UTC

Amount Due: \$0.66 USD

This statement is generated automatically on demand. All figures are based on measured benchmarks from production runs.