

Systems Status Report — 2026-04-27

Systems Status Report

Generated: Monday, April 27, 2026 — 15:18 UTC
Prepared by: Sam (👤)
Scope: All managed infrastructure — Gateway host, Node2, TS (Proxmox), and services

1. Gateway Host (claw) — 10.50.15.226

Metric	Value
Uptime	5 days, 14 hours
Kernel	6.8.0-110-generic
CPU	Intel Xeon E5-2680 v4 @ 2.40GHz (8 cores / 8 threads)
RAM	4.2 GiB used / 7.8 GiB total (54%)
Swap	768 MiB used / 4.0 GiB total
Root Disk	35 GiB used / 124 GiB (29%)
Load Average	1.16 / 0.84 / 0.76
Network	ens18: 10.50.15.226, Tailscale: 100.107.29.33
Gateway	Running (PID 3256056), port 18789, loopback-only
Sessions	94 active
Tasks	1 running, 27 issues, 2407 total background tasks

OpenClaw Status

- **Heartbeat:** 30 min (main)
- **Default model:** ollama/qwen3.5:cloud (262k ctx)
- **Runtime model:** node2/llamacpp
- **Cache:** On (1215 entries)

Recent Errors (last 6 hours)

- ⚠ **LLM timeout** — node2/llamacpp timed out at 14:55 UTC (125s), fell back to claude-opus-4.6
- ⚠ **Copilot auth failure** — HTTP 403 token exchange at 14:55 UTC
- ⚠ **503 Loading model** — node2/llamacpp returned 503 at 15:07 UTC (model loading)
- ⚠ **Memory sync failed** — CredentialsProviderError at 15:10 UTC

Security Audit (3 CRITICAL)

1. **CRITICAL:** Exec security=full is configured for main agent — should use allowlist
2. **CRITICAL:** Open channels can reach exec-enabled agents (matrix groupPolicy=open)
3. **CRITICAL:** Open groupPolicy with elevated tools enabled — prompt injection risk

2. Node2 (GPU Server) — 172.16.254.101

Metric	Value
Uptime	4 days, 12 hours
RAM	5.9 GiB used / 62 GiB total (10%)
Swap	1.8 MiB used / 8.0 GiB total
Root Disk	174 GiB used / 249 GiB (74%)
Load Average	0.51 / 1.65 / 3.13
Tailscale	100.91.255.3

GPU Status — 2× NVIDIA Quadro P6000 (24 GB each)

GPU	Temp	Power	Limit	Util%	VRAM Used	VRAM Free	Driver
0	54°C	69.2 W	250 W	0%	21,361 MiB (87%)	3,077 MiB (13%)	580.126.20
1	45°C	65.3 W	250 W	0%	21,865 MiB (89%)	2,573 MiB (11%)	580.126.20

Power state: D0 (active) on both GPUs
CUDA: 13.0
Avg power draw: ~66.8 W (over 2.37s window)

llama-server.service

- **Status:** □ Active (running since 15:07 UTC)
- **PID:** 3388390
- **Model:** llamacpp.gguf (103 GiB model directory)
- **Port:** 11434
- **GPU layers:** all (both GPUs)
- **Threads:** 12 / 12 batch
- **Context:** 262,144 tokens
- **Concurrency:** -np 1 (single-slot — bottleneck confirmed)
- **Memory:** 5.2 GiB RSS
- **Recent activity:** Processing requests normally (task 2432 at 15:15 UTC)

Disk Breakdown (largest)

- /opt/models/gguf — 103 GiB
 - /home/localadmin/projects — 4.6 GiB
 - /home/localadmin/logs — 73 MiB
 - /home/localadmin/tmp — 43 MiB
-

3. TS (Proxmox ESXi) — 172.16.254.5

Metric	Value
Uptime	9 days, 11 hours
RAM	90 GiB used / 251 GiB total (36%)
Swap	None
Root Disk	3.5 GiB used / 897 GiB (1%)
Load Average	0.57 / 1.60 / 3.27

VM Status

VMID	Name	Status	RAM	Bootdisk
100	NODE1	❑ STOPPED	65 GB	128 GB
101	NODE2	❑ Running	65 GB	256 GB
200	fs.9xc.local	❑ Running	4 GB	32 GB

4. Services Status

Service	Status	Notes
OpenClaw Gateway	❑ Running	PID 3256056, port 18789
llama-server (node2)	❑ Running	Port 11434, processing requests
n8n	❑ DOWN	No systemd unit, no binary (only config at ~/.n8n/)
gog-bridge	❑ DOWN	No service file, no binary
Matrix (openclaw-matrix)	❑ DOWN	Service unit not found

5. Cron Jobs (10 total)

Name	Schedule	Status	Last Run
node2-metrics-scraper	Every 5 min	❑ OK	3 min ago
node2-p6000-monitor	Every 4h	❑ OK	1h ago
Daily Systems Status	Weekdays 18:00 UTC	❑ OK	3 days ago
daily-systems-status (main)	Every 1d	❑ OK	9h ago
Memory Monitor	Daily 08:00 UTC	❑ OK	7h ago
Daily Node2 Cost Report	Daily 12:00 UTC	❑ OK	3h ago
6AM CST Morning News	Daily 12:00 UTC	❑ OK	3h ago
node2-metrics-aggregator	Every 1d	❑ OK	2h ago
Daily Self-Improvement	Daily 14:00 UTC	ERROR	1h ago
Monthly Invoice - Ollama	1st of month 09:00 UTC	❑ Idle	—

Health: 8/10 OK, 1 ERROR, 1 IDLE

6. Outstanding Issues

🔴 Critical

1. **Security: Exec full trust + open groupPolicy** — 3 critical security findings in openclaw status. Recommend tightening to allowlist with ask prompts.
2. **Daily Self-Improvement cron** — Failed at 14:00 UTC today.

🟡 Medium

3. **n8n + gog-bridge dead** — No binaries or service files. Config/DB still intact at ~/.n8n/. Needs full reinstall.
4. **Matrix service** — Unit file missing.
5. **node2 root disk at 74%** — 174/249 GiB used. Model directory (103 GiB) is the largest consumer.
6. **llama-server -np 1** — Still single-slot (11 days unresolved). ~700 MiB free per GPU available for a 2nd slot.

🟢 Low

7. **NODE1 VM stopped** on Proxmox — intentional?
 8. **Copilot auth (HTTP 403)** — intermittent auth failures in gateway logs.
 9. **Model loading 503s** — node2/llamacpp occasionally returns 503 during model load.
-

7. Summary

Overall status: ⚠️ **DEGRADED**

- 3 of 4 hosts are operational and stable
 - Gateway host running normally with 3 critical security findings to address
 - Node2 GPU server healthy — both GPUs idle but VRAM-pinned, llama-server processing
 - TS Proxmox stable — 2/3 VMs running
 - **n8n and gog-bridge are down** — automation pipeline broken
 - Self-Improvement cron job errored today
 - Security posture needs tightening (full exec trust + open groupPolicy)
-

Report generated by Sam 🐱 — OpenClaw Personal Assistant