

Systems Status Report

Report Date: 2026-05-29 23:36 UTC Generated by: Sam (OpenClaw Ops Butler)

Host: claw (Gateway)

Metric	Value	Status
Uptime	4 days, 8 hours	Healthy
CPU Load	0.39 / 0.32 / 0.28	
RAM	2.7G used / 7.8G total (5.0G available)	
Disk (/)	55G used / 124G (47%)	
Swap	512K used / 4.0G	All active
Services	openclaw-gateway, n8n, gog-bridge	
Tailscale	Connected (100.107.29.33)	
Network	ens18: 172.16.254.226/24	Clean
Errors (24h)	None	

GPU / Inference Nodes

node1 (2× Quadro P6000 — 48GB) — **CRITICAL DISK**

Metric	Value	Status
Uptime	4 days, 8 hours	Idle
CPU Load	0.00 / 0.01 / 0.11	
RAM	2.6G used / 46G (44G available)	CRITICAL
Disk (/)	123G used / 123G (100%)	
Llama server	HTTP 200 on :11434	
Model	Qwen3.6-35B Q4_K_M	

ACTION REQUIRED: Disk is completely full. Any file write will fail. Llama inference is still running but no new processes can write to disk. Immediate cleanup needed.

node2 (1× Quadro P6000 — 24GB) — **ALL GREEN**

Metric	Value	Status
Uptime	4 days, 8 hours	Moderate
CPU Load	0.89 / 0.41 / 0.25	
RAM	2.6G used / 46G (44G available)	
Disk (/)	182G used / 501G (38%)	
Llama server	HTTP 200 on :11434	
Model	Qwen3.6-35B Q4_K_M	

pve3090-111 (2× RTX 3090 — 48GB) — ALL GREEN

Metric	Value	Status
Uptime	4 days, 19 hours	Idle
CPU Load	0.07 / 0.04 / 0.00	
RAM	23G used / 62G (39G available)	
Disk (/)	24G used / 63G (41%)	
Llama server	HTTP 200 on :11434	
Model	Qwen3.6-35B-A3B Q8_K_XL	

hermes (CPU-only — NousResearch) — ALL GREEN

Metric	Value	Status
Uptime	4 days, 8 hours	Idle
CPU Load	0.00 / 0.00 / 0.00	
RAM	700M used / 7.8G (7.1G available)	
Disk (/)	8.6G used / 61G (15%)	
SSH	10.50.15.231	

Proxmox / Infrastructure

TS.9xc.local (Supermicro SYS-6028R-T — 2× Xeon E5-2680 v4, 251GB RAM) — UP

Metric	Value	Status
Uptime	4 days, 8 hours	Heavy
CPU Load	5.96 / 5.54 / 5.34 (7 sockets, 112 threads)	
RAM	148G used / 251G (103G available)	
OS	Proxmox / Debian Trixie	

ZFS Pools:

Pool	Size	Used	Free	Usage	Status
rpool	928G	9.2G	919G	1%	Healthy
fast-store	3.62T	1.18T	2.45T	32%	Healthy

Note: Load average ~5.5 is normal for a 7-socket 56-core host. CPU utilization is healthy at ~10%.

USM1 (Supermicro — Xeon E5-2620 v4, 62GB RAM) — UP

Metric	Value	Status
Uptime	2 days, 22 hours	
CPU Load	0.17 / 0.42 / 0.52	Light
RAM	42G used / 62G (19G available)	
Disk (rpool)	14G used / 346G (5%)	
GPU	None	CPU-only
OS	Proxmox / Debian Trixie 13.4	

LLM Inference Summary

Provider	Host	GPU	Model	Status	HTTP
node1	172.16.254.100	2× P6000 (48GB)	Qwen3.6-35B Q4_K_M		200
node2	172.16.254.101	1× P6000 (24GB)	Qwen3.6-35B Q4_K_M		200
pve3090-111	pve3090-111	2× RTX 3090 (48GB)	Qwen3.6-35B- A3B Q8_K_XL		200
hermes	10.50.15.231	None	—		UP

Model Routing: - Main session: pve3090-111/llamacpp (VM111) - Sub-agents: node1 OR node2 (50/50 split)

N8N Automation

Metric	Value	Status
Process	Active (npx, PID 806)	
Port	:5678	
API endpoint	/api/v1/	
Gog bridge	Active on :18790	

OpenClaw Cron Jobs

Job	Schedule	Last Status	Next Run	Notes
Server CPU Temp Check	Every 4h	OK	Tonight 00:00	Staggered +1min

Job	Schedule	Last Status	Next Run	Notes
Obsidian Vault Daily	Daily 12:00 UTC	OK	Tomorrow 12:00	
Morning News Briefing	Daily 12:00 UTC	ERROR (×2)	Tomorrow 12:00	LLM network error
Daily Vault Git Push	Daily 13:00 UTC	OK	Tomorrow 13:00	
Vault Weekly Summary	Sunday 13:00 UTC	ERROR (×1)	Sunday 13:00	Timeout on model call
Monthly Invoice	1st of month 09:00 UTC	OK (delivered)	June 1	

Total jobs: 6 (4 healthy, 2 with recent errors)

Issues & Action Items

CRITICAL — node1 Disk Full (100%)

- **Impact:** No writes possible. Llama server still running but no new processes.
- **Recommendation:** Clean up old model files, build artifacts, or logs on node1.
- **Affected services:** Llama inference (node1/llamacpp), any file writes.

Morning News Briefing — Consecutive Failures

- **Issue:** 2 consecutive failures due to LLM network error (cloud model endpoint).
- **Recommendation:** Check network connectivity from isolated sessions. May be transient.
- **Risk:** Missed daily briefings if not resolved.

Vault Weekly Summary — Timeout

- **Issue:** Job timed out during model call (301s duration limit).
- **Recommendation:** Increase timeout or simplify weekly summary scope.

Overall Health Summary

Category	Status
Gateway (claw)	Healthy
GPU Nodes	node1 disk critical
Infrastructure (TS, USM1)	Healthy
LLM Inference	All responding
N8N / Automation	Running
Cron Jobs	2 with recent errors

System-wide health: **CAUTION** — One critical disk issue (node1), two cron errors (likely transient). Everything else operating normally.

Report generated by Sam — OpenClaw Operations Butler Next scheduled report: daily at 12:00 UTC