

VM111 Systems Status Report

Generated: 2026-04-27 16:47 UTC | Host: vm111 (user1@vm111) | Report by: Sam

System Overview

Parameter	Value
OS	Linux 6.8.0-110-generic (Ubuntu) — x86_64
Uptime	1h 52m
Load Average	0.36 / 0.21 / 0.28 (16 cores — healthy)
Users logged in	5

CPU

Parameter	Value
Model	Intel Xeon E5-2680 v4 @ 2.40GHz
Sockets	2
Cores per socket	8
Threads per core	1 (no HT)
Total logical CPUs	16

Memory

Type	Total	Used	Free	Available	Status
RAM	62 GiB	6.8 GiB (11%)	28 GiB	55 GiB	OK
Swap	8 GiB	2.2 MiB	8 GiB	—	OK

Disk

Filesystem	Type	Size	Used	Available	Use%	Status
/dev/mapper/ubuntu--vg-ubuntu--lv	ext4	492G	116G	351G	25%	OK

GPU — 2× NVIDIA GeForce RTX 3090

Parameter	GPU 0	GPU 1
Driver	580.126.20 (CUDA 13.0)	
Temp	45°C	55°C
Power	151W / 350W	163W / 350W
VRAM Used	22,012 MiB / 24,576 MiB (89.6%)	22,930 MiB / 24,576 MiB (93.3%)
GPU Util	0%	0%
Fan	33%	33%
Process	llama-server (PID 1683) — model loaded across both GPUs (ts 10,12)	

Status: VRAM heavily allocated (~45 GiB / 48 GiB), but GPU utilization at 0% — model loaded, idle between requests.

Network

Interface	State	Address
lo	UNKNOWN	127.0.0.1/8
enp6s18	UP	192.168.68.111/22
tailscale0	UNKNOWN	100.123.212.45/32

Services (Running)

Service	Status	Description
llama-server.service	active	llama.cpp Server (CUDA - 2x RTX 3090)
sshd.service	active	OpenSSH Server
tailscaled.service	active	Tailscale node agent
nvidia-persistenced.service	active	NVIDIA Persistence Daemon
systemd-networkd.service	active	Network Configuration
systemd-resolved.service	active	Network Name Resolution
systemd-timesyncd.service	active	Time Synchronization
dbus.service	active	D-Bus System Message Bus
multipathd.service	active	Device-Mapper Multipath

Top Processes by CPU

PID	User	%CPU	%MEM	RSS	Command
1683	llamacpp	33.0	10.7	6.8 GB	llama-server (model: llamacpp.gguf + mmproj)
1750	user1	2.2	0.0	24 MB	nvtop
1752	root	1.2	0.0	10 MB	sshd (current session)
864	root	0.7	0.0	45 MB	tailscaled

llama-server Configuration

Parameter	Value
Model	/opt/models/gguf/llamacpp.gguf
Multi-modal projector	/opt/models/gguf/mmproj.gguf
Listen	0.0.0.0:11434
Context size	262,144 tokens
GPU layers	all (offloaded)
Batch size	2024 (ub: 1024)
Threads	12 (batch: 12)
GPU split	ts 10,12 (layers per GPU)
Parallel	-np 1 (single slot — bottleneck)
Flash attention	enabled
KV cache	q8_0 (K+V), 16 GB RAM fallback
Temp / Top-P / Top-K / Min-P	0.6 / 0.95 / 20 / 0.00

Summary & Observations

Area	Status	Notes
CPU / Load	Healthy	0.36 load on 16 cores. Plenty of headroom.
Memory	Healthy	11% RAM used (6.8/62 GB). 55 GB available.
Disk	Healthy	25% used (116/492 GB). No concern.
GPU VRAM	⚠ High	89-93% VRAM on both GPUs. Model loaded, idle. Expected for current config.
GPU Temp	Normal	45°C and 55°C — well within limits.
Network	Healthy	LAN (192.168.68.111) + Tailscale (100.123.212.45) both up.
Services	All Running	16 services active, no failures.
llama-server	⚠ -np 1	Still single-slot. Only one concurrent inference stream. Consider -np 2 for dual sessions.

